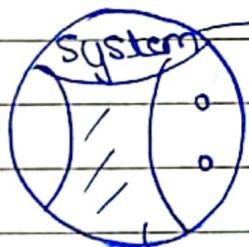


DATABASE

imp

- File organization didn't have any data abstraction, so DBMS.
- DBMS: concept / hypothesis to facilitate store, manage & access data.
- RDBMS is the implemented version [External view, Logical view, Physical view].
- Any piece of data is physically stored in files only, but logically in the form of tables, in dbms.



System.dbf contains the metadata; list of objects, list of col's, location of data.

table Space

* Our SQL query (ext. view), hits system.dbf, → logical view → Physical view extracts data out of file blocks & back to user. [Vimp]

INTEGRITY CONSTRAINTS

- Data integrity constraints

unique
not Null
check
primary key

validate correctness accor
to value.

→ Unique constraint can't validate whether the value is NULL/not.

*FUNCTIONAL DEPENDENCIES

- determinant → dependent
 - ▷ det. gets duplicated, dep. also gets duplicated
 - ▷ det. gets unique, dep. can have any value

- (ii) $R \rightarrow R$ always valid [Tautology]
 - (iii) $A \rightarrow BC \equiv A \rightarrow B \ \& \ A \rightarrow C$.
 - (iv) $A \rightarrow B \ \& \ B \rightarrow C$ valid -
 $\therefore A \rightarrow C$ also valid -
 - (v) $\because BC \rightarrow BC ; \quad \therefore BC \rightarrow B \ \& \ BC \rightarrow C$ also
valid.

leg → Pg 7 [Instance of relation] ✓

- 4) Super Keys: if all attributes of relation can be derived by them.
 can be single / composite.

5) Candidate Keys: if AB & A can derive all attributes, ACAB

$AB \rightarrow \text{SuperKey}$; $A \rightarrow CK$.
 ∵ CK is minimal subset of SK.

1 * Normalization: Process applied on relation to reduce deg of redundancy.
→ Data redundancy causes:
mem waste, updtng/insertn/deln among
eg → pg 11 [redundancy] ✓

* NORMAL FORM: indicates amt of redundancy present in it. (imp)

Normal form
NF & 1
deg of redundancy

- Step 1: find out highest NF
- Step 2: decompose into higher ~~NF~~

Step 1: Find out highest NF:

Step(1) Find out all CKS

~~ex~~ → Pg 13 (Sne, Pne)

Step(2) Find prime / Non-Prime attributes.

[eg] → Pg 6 (Rev Copy) ✓ Imp

* Problem with redundancy:
eg → Pg 15 [T/U/D anomaly] ✓

Step(3) Identify dependencies

(ii) Full Dependency

CK/SK → Prime / Non-Prime

- never comes redundantly. ✓

RHS = Null
(PD & TD)

Date
Page 4

(ii) Partial dependency



NOTE: If simple CK, no PD exists

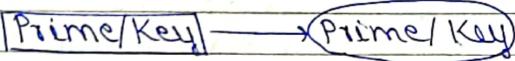
Imp

(iii) Transitive dependency



NOTE: If relation has only Prime Attrib.
(PD & TD = 0) ✓ Imp

(iv) Overlapping CK dependency



Step 4) Types of Normal Forms

(i) 1NF: all values in relation must be atomic. Every relation → 1NF

(ii) 2NF: No PD

(iii) 3NF: No PD, no TD

(iv) BCNF: only FD (no OCD)

✓

* Hierarchy of NF



weakest

Vimp

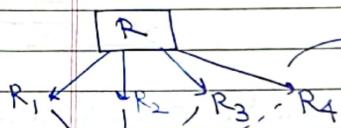
strongest

[eg] → Pg 22 [Q2] (OCD) ✗

[eg] → Pg 23 [Q3] (r. K defin) ✗

[eg] → Pg 25 [choose best clause & comb. missing attribute] ✗

STEP - 2: Decompose into higher NF



Properties of Normal:
(i) lossless decompos.
(ii) Dependency preserv.

↓ ↓ ↓ ↓
R → guar. availib. of data

(i) lossless decomposition: By join R, R₁, R₂, R₃, R₄ if we can derive all data in eg. R.

(ii) dependency preserving: By using the projected FDs, if we can derive all FDs. Guarantees integrity constraints.

[eg] → Pg 29 [1NF → 3NF form] ✗

Imp [normal
normalization]

1*) Procedure to find decomposition is lossless or not.

$$[Y \rightarrow Y \checkmark \quad Y \rightarrow X \checkmark \quad X \rightarrow Y \times]$$

eg → Pg 11 [Rev Copy] (i) entire row filled with X, lossless decamp.

2*) Procedure to find decomposition is dependency preserving or not.

- (i) Find all closures.
- (ii) Apply the FD projections on decomposition.
- (iii) Check if you can derive all FDs of relation [using FD projections].

eg → Pg 12 \times

4*) BCNF Algorithm

Guarantees lossless decomposition, but not dependency preservation.

[may/may not]

eg → Pg 34 \times $[X \rightarrow Y \text{ or } (\text{not FD})]$

$$XY \rightarrow R - f[Y]$$

→ Covers F if all the FDs of F are implied by G.

5*) 3NF Decomposition Algorithm

Guarantees lossless + dependency preserving decomposition.

\times

Date
Page
6

imp [Other Method also]

classmate
Date
Page
 \neq

• Minimal cover: eliminates unnecessary attributes on determinant side.

$$AB \rightarrow C, A\bar{B} \rightarrow C$$

$$A \rightarrow C$$

$$A \rightarrow B$$

imp

• Canonical cover: on dependent side.

$$AB \rightarrow CD$$

$$A \rightarrow C$$

eg → Pg 39 [R(A, B, C, D, E)] \times

② RELATIONAL ALGEBRA

(i) Set Operations: { U, ∩, −, × }

(ii) Native Operations: { σ, π, ⋈, ÷, ⋉ }

• degree: No. of attributes/columns

• Cardinality: No. of rows

• domain: a char, b int \times

imp

* To be union compatible, degree & domain of respective attributes must be same.

(i) Set Theoretic Operators

(i) Union:

$$d(R \cup S) = d(R) = d(S)$$

$$\max(|R|, |S|) \leq |R \cup S| \leq |R| + |S|$$

(ii) Intersection:

$$d(R \cap S) = d(R) = d(S)$$

$$0 \leq |R \cap S| \leq \min(|R|, |S|)$$

imp

(iii) set-diff: $d(R - S) = d(R) = d(S)$
 $0 \leq |R - S| \leq |R|$

(iv) Cross-Product: $d(R \times S) = d(R) + d(S)$
 $|R \times S| = |R| * |S|$
 ↳ result need not be union compatible.

(ii) NATIVE OPERATIONS [like relations]

(1) Join: subset of Cartesian Product.
 ↳ Join Type: inner / outer / left outer / right outer

↳ Join Condition: natural / ON / using.

a) Join Condition
 → Natural: at least 1 common column only = used, eliminates duplicates of common columns [Equi Join]
 → ON card / Theta Join: no prerequisites, any operator used, duplicates not removed.
 → Using: same / all col used.

If all common columns are used, result will be like Natural Join.

b) Join Type:
 → Inner: allows only matched rows.
 → LOJ: Matched + Left's Unmatched
 → ROJ: Matched + Right's Unmatched.

Data Page 8

classmate
Data Page 9

[eg] → Pg 48 ✓

Join considers a row as qualified if it gets matched atleast once.

(2) DIVIDE:

For R : S: Att. of S \subseteq Att. of R

Divide considers a row as qualified if it gets matched with all the ref. values.

[eg] → Pg 49 ✓

(3) Selection (σ): Produce logical horizontal subset of rows.
 ↳ where clause in SQL.

$$d(\sigma_{\text{result}}) = d(\text{original})$$

$$|\sigma_{\text{result}}| \leq |\text{original}|$$

(4) Projection (π): Produce vertical subset of rows by eliminating duplicates.
 ↳ 'Select with Distinct' clause in SQL

$$d(\pi_{\text{result}}) \leq d(\text{original})$$

$$|\pi_{\text{result}}| \leq |\text{original}|$$

(5) Rename: doesn't change deg / cardinality.

③ SQL

- Ansi Stand SQL

- DDL: create, alter, rename, truncate, drop
- DML: Insert, upd, del, sel
- DCL: grant, revoke

TCL: commit, rollback

DDL → truncate: only user data gets deleted
drop: User + Metadata gets deleted

* Truncate: delete all rows of Table
delete: delete 1/some/all rows

* SELECT (DML Command)

eg → Pg 52 [order of exec] OR AND NOT

eg → Pg 53 [operator precedence]

where deptno = 10 where deptno
or deptno = 20 ≈ IN (10, 20, 30)
or deptno = 30 ✗

imp → * sal <= 5000 and ≈ sal BETWEEN
sal >= 2500 2500 and 5000

* LIKE: where ename like '%',
where name like 'S%'

eg → Pg 55 [Exc]

* FUNCTIONS IN SQL

Single Row Funct

→ Applied on 1 row at a time to get one result per row.



Multirow / Group Aggregate funct

[max(), min(), sum(), avg(), count()]

can't be used in where clause

imp can be used with having

* Group By

diag → Pg 22 [Rev Copy] (6 cases with Group By clause → imp) ✗

NOTE: All the non-group exp that exists in the Select clause along with group exp must be present in the group by clause. imp

* Order By → sorting by an attrib.

→ By default, SQL displays the result set in the same behaviour/order as data got inserted.

eg → Pg 60 [many exp in order by clause]

* Subqueries

→ Single Row SQ

→ Named SQ → Multi Row SQ

→ Correlated SQ ✓

1) (i) Single Row Sub Query

eg → Pg 62 [Q1] ✓ → SQL returns 1 Value

(ii) Multi Row Sub Query

eg → Pg 24 [Rev Copy] ✓ (Imp)

Normal SQLs get executed only once irrespective of no. of rows we have in Table ✓

2) Correlated Sub Queries

- SQL gets executed as many times as no. of rows in Table.

eg → Pg 66 [Imp] ✓ (Imp)

* SET OPERATIONS IN SQL

• Union in SQL eliminates duplicates

eg → Pg 66 ✓ (Imp)

where clause maps row by row ✓

* PREVIOUS YEARS

Q) If value of a relation is given, we can say A → B (may/may not)

CLASSMATE
Date _____
Page 12

CLASSMATE
Date _____
Page 13

→ We need full picture.

2) Any relation with 2 attributes is always in BCNF (Binary Rel).

eg → Q2-40 [Max no. of SK] ✓ [ME]

eg → Q2-24, Q2-32 [Max Tuples when NJ is applied] min(IRI, ISI) ✓

* SQL query will work even if there are no indexes on relations. SQL doesn't permit attribute names to be repeated in same relation. (Imp)

* To check for lossless decomposition:

1. $AT(R_1) \cup AT(R_2) = AT(R)$.
2. $AT(R_1) \cap AT(R_2) \neq \emptyset$.
3. $AT(R_1) \cap AT(R_2) \Rightarrow$ must be a CK for a relation ✓ (Imp)

④ TRANSACTION MANAGEMENT

Transaction: group of DML inst. to convert data from 1CS to another. It must satisfy ACID Properties:

(i) Atomicity: either all instructions a) trans. must get executed / None embedded b/w [Begin trans, end trans]

maintained
help of shadow
paging

Imp

Data
Page 14.

(ii) Consistency: data b4 trans. & after
trans. must be consistent (cannot
→ Shadow Paging used)
↳ copy of data item before &
after transaction is kept.

TCL

(1) commit: all modif performed by
trans is transferred from PB cache
to SM.
(2) rollback: ignore all modif's &
replace such value with previous
committed values.

(iii) Isolation: a trans can't affect any
other trans during execution

Schedule → serial

→ concurrent

→ Serial Schedule

all unit of trans get executed non-
preemptively.

If n Transaction = n!, Serial Sch

Adv: Guarantees data consistency

Disadv: High WT, low throughput

→ Concurrent Schedule

Adv: low WT

Disadv: can't guarantee consistency

* Conflicting Operations ↴ due to

Data
Page 14.

classmate

Date
Page 15.

diff trans / same data item
W-W : lost update Problem
W-R : dirty read Problem
R-W : unpredictable read prob

* Serializable Schedule

A concurrent Schedule is serializable if
result = result of any of Serial S

[eg] → Pg 75 [Wh conflict] ✓
(Common DB cache, local cache)

[eg] → Pg 30 [RCV] Pg 78 (Imp Q) ✓

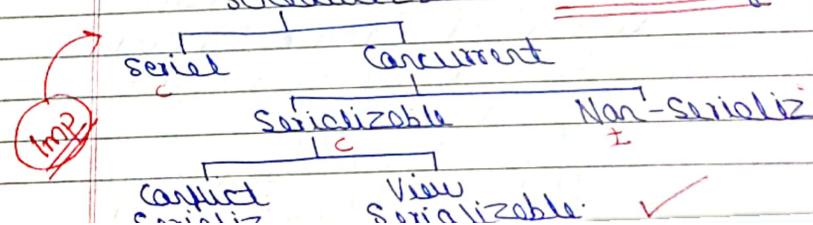
conflicting operations → Possibility of data
operations inconsistency not
always ✓

if serializable → data is consistent

⇒ Conflict Serializable: If given conc
schedule is conflict equivalent to
one of the serial schedules.

[Same for view Serializable]

Schedule [based on consistency]



Gate

* If Schedule is either conflict / view
(both \Rightarrow Serializable)

1) Conflict Serializable

\rightarrow 2 Schedules are conflict equivalent
if order of conflicting oper is same

eg Pg 81 ✓

To find conflict Serializable
use Precedence Graph
Alg.

Gate

- Perform Topological ordering on
Graph to find equivalent Serial Sched.
for concurrent Schedule

eg Pg 84 [Gate 2015] ✓

eg Pg 86 ✓

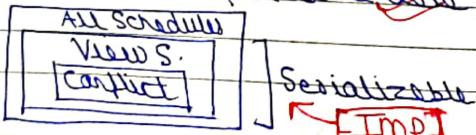
2) View Serializable

Imp

(i) Trans. that perform init read on
S1, should also perform init read
on S2, also on same data item.

(ii) Same with final write.

(iii) If RWConflict exist on same data
item in S1, same RWConflict must
exist in S2 also, on same data item



Date
Page 16

classmate
Date
Page 17

eg Pg 89 [Q1] (View Serial) ✓

eg Pg 90 ✓ If there exist more than 1
trans. with init read, & atleast 2
perform write on that data item not VS
init read = write pehle na write.
final write = write badal na write.

(iv) Durability: All modif. performed by
committed trans. are assumed to be
present in SM (no matter what type
of failure). Gate ✓

* WAL Protocol: All modif. by every
commit. trans. must be written to a
redo log file before actual data file.

used in:

- Analysis Phase
- Redo Phase
- Undo Phase.

* RECOVERY

→ redoing all entries of
redo log file ✓

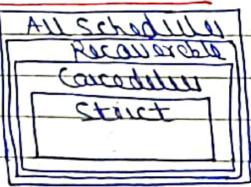
eg Pg 93 [Diag] Gate ✓

eg Pg 94 [Diff b/w commit & Check
Point] ✓

Imp

* Schedules based on Recoverability

- (i) Recoverable Schedule: if $T_i \rightarrow T_j$ exists, $T_i = \text{write}$, $T_j = \text{read}$; then completion order must also be T_i, T_j .
(imp)
- (ii) Cascading rollbacks are caused.
- (iii) Cascading Schedule: dirty read not allowed. Trans can read or data items update by only committed transaction.
- (iv) Strict Schedule: No dirty / last update; Trans can't read / write on the [modif'd] done by other trans unless committed.



eg → Pg 99 ✓

eg → Pg 100 / Pg 101 ✓

Find out highest recoverable Schedule.
(imp)

* CONCURRENCY CONTROL

- RDBMS uses locking mechanism.

Date _____
Page _____ 18

CLASSEEE
Date _____
Page _____ 19

Lock Manager accept/reject lock req. of transactions using Lock compatibility Matrix.

if trans. wants to perform
read : shared
read / write : exclusive ✓

Concurrency control aims for producing Strict Serializable Schedule always.

eg → Pg 36 [Rev Copy] Lock compatibility Matrix. ✓

As per RDBMS, Locking & unlocking should be performed as per a protocol called 'Locking Protocol'.

1 Two Phase Locking Protocol

- Graving Phase: Trans only allowed to acquire locks.

- Shaking Phase: Trans only allowed to release acquired locks.

Problem: (i) deadlock Problem.

(ii) Cascading rollback Prob.

(not strict)

→ always uses Serializable Schedule

eg → Pg 104 (Deadlock Problem) ✓

2 Strict 2 Phase Locking Protocol

A trans. can release all the acquired exclusive locks, fill it commits
 Disadv: deadlock Problem
 Adv: Produce strict Serializable Schedule ✓

3) Time Stamp Based Protocol

- Unique TS will be assigned to each & every transaction
 Adv: No deadlock. ✗ Gate

4) ER DIAGRAMS

* Relational integrity constraints

* FK : PK/CK of master table that exists in child table.
 → A value is allowed in the FK column if it has reference to one of the values of PK column of parent table.

Relations can have many CK, but out of them, only 1 becomes PK ✗

must be defined
 → on delete : If PK value gets deleted, cascade all other FK values also deleted
 , on delete : User manually updates No action FK values, so that if PK deleted, FK no action value

* Basic ER Diagrams

combine child + relationship to min. no. of tables. many side

eg → Pg 40 (diag) [Rev Copy] ✗

* Participants

Complete / Incomplete /
Total Partial

- FK values must be NULL in child Table (By default) [Double line]
- FK values can be NULL [Dash line]

* Weak / Strong entities

→ doesn't have PK

→ rel. with SE, with full participation
col list of col + PK of SE of WE → PK

eg → Pg 114 [diag] Imp ✗

* Gate Q.S.

eg → Pg 115 (2008) [ER diagram min] ✗

eg → Pg 119 (2005) [Completed SQ] ✗

eg → Pg 121 [NJ, Union, Intersection
difference, divide] - RA ✓

eg → Pg 123 [RA & TC] ✗ (imp)

* Question on Normalization

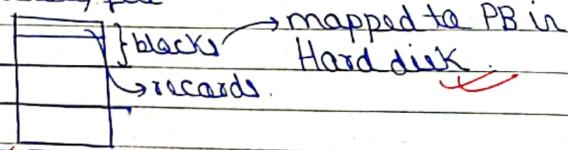
eg → Pg 42 (Rev Copy) ✗ (imp)

eg → Pg 43 [R(A,B,C,D,E,F,G)] ✗

⑥ FILE ORGANIZATION

In OS, we needed to access file from begin to end, but in DBMS, we need some record from a file. ∵ file organization also diff.

table/file



* Strategies for storing file of records with blocks:

- (1) Spanned: allows partial part of record to be stored in block. No waste of mem - variable length records. (imp)
- (2) Unspanned: No record can be stored

Date Pg. 22

CLASSmate

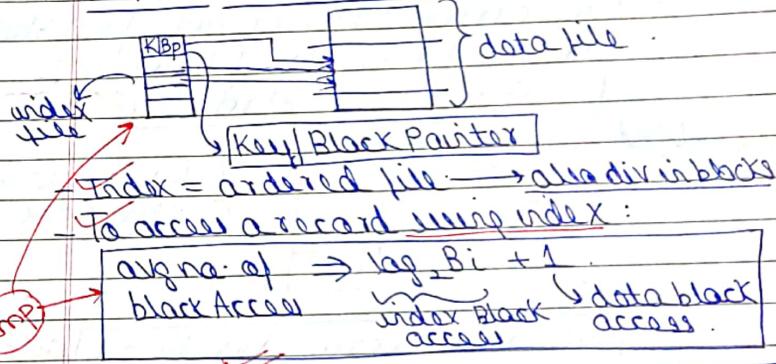
Date Pg. 23

in more than 1 block.
fixed length record. ✗

* Organization of records in a file

- (i) Ordered file organization: records are ordered, based on a field. Binary Search used to search record. Inorder → expensive
- (ii) Unordered file organization: records in file are inserted wherever space is available; at EOF. (Linear Search)

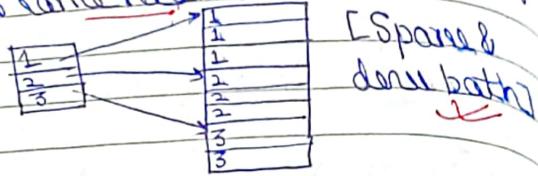
* INTRO TO INDEXING



* Classification of Indexing: (imp)

- (1) Dense index: if index entry is created for every distinct search key value.

(2) Sparse index: index created only for some records.



* TYPES OF INDEXING

Single level index

Multi level index
[B / B+ Trees]

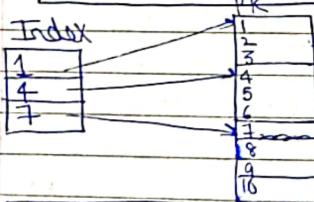
↓ PI, CI, SI



* Primary Index [PK + Ordered]

- Index created for first record of each block

$$\text{No. of index entries} = \text{No. of data blocks}$$



{may/maynot}
Sparse index

$$\text{Avg Block Access} = \log_2 B_i + 1$$

eg → Pg 7 [Primary indexing] ✓ (imp)

* Clustering index (Non Key + ordered)

- ordered or non-key. : repeated.

→ for every distinct key value : 1 index record. [Sparse + Dense]
: Avg Block Access = $\log_2 B_i + 1$

* Secondary index [Non Key / CK + ordered]

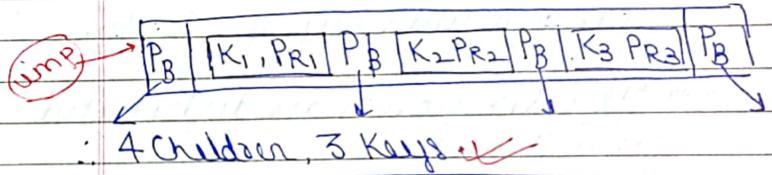
→ Every record
index created for every record in file
[Dense, not Sparse] (imp)
- ST provides a secondary means of
accessing file for which PK already
exists.

eg → Pg 10 [Secondary Index] (Gate)

✓ An index is clustered if data records
of file are organised in the same
order as the data entries of index.

* B-TREES

→ Generalization of Multi level index.
→ Node / Block Pointer } + Key (at each
level).
→ record pointer



$$P = \text{order}/\text{max children node can have}$$

(i) Root: $\text{children} \in [2, p]$
 $\text{Keys} \in [1, p-1]$

(ii) Internal node:
 children: $\lceil \frac{p}{2} \rceil \text{ to } p$
 Keys: $\lceil \frac{p}{2} \rceil - 1 \text{ to } p-1$

Same for leaf node

eg → Pg 14 (max children/order = 23) ✓
eg → Pg 16 [VimpQ] (4 levels) ✓

* UNDERFLOW / OVERFLOW in B-Tree

- overflow: if keys exceed ' $p-1$ '
underflow if keys < $\lceil \frac{p}{2} \rceil - 1$

Inversion problem Deletion issue

~~1)~~ B-Tree Search

- first get into root; if match: Stop
[compare with key, if $<$; then go left]
 - All index records are distributed over B-Tree.
 - Search the node, till we find a key greater than the key we are searching for, & go left

- Split left node key into 2 parts & promote median key to Parent.
→ Pg. 19 [Eg. 1st Inversion] ✓

\Rightarrow When we Split the node;

- $p = \text{add}$ $p = \text{even}$ $\left\lfloor \frac{p-1}{2} \right\rfloor, \left\lfloor \frac{p-1}{2} \right\rfloor + 1$ $\left[\text{but it is small} \right]$

Split when overflow occurs

B-Tree deletion

$$\text{order} = 5$$

Borrow:

Merge

```

graph LR
    CG1[CG] --> AB1[AB]
    CG1 --> DE1((DE))
    DE1 --> G2[G]
    CG1 --> HI1[HI]
    G2 --> AB2[AB]
    G2 --> DE2[DE]
    G2 --> HI2[HI]
  
```

~~(3) B+ TREES~~

• B-Trees: distribute (Key-Record p) throughout the tree

→ B+Tree:

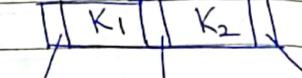
Internal node: remove record pair
only Key + Node Pointer

(imp)

: Breadth ↑, depth ↓, Time ↓

Leaf node: (Key, Rp) pair + 1 NP.

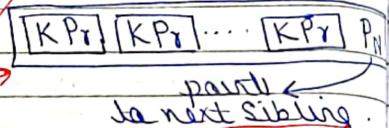
* Internal Node



children: $\lceil P/2 \rceil$ tap p

Keys : $\lceil P/2 \rceil + 1$ tap p - 1

* Leaf Node



pair ←
to next sibling.

(imp)

* Order of Non-leaf: max no. of children it contains.

* Order of Leaf: max no. of K-V pairs present inside leaf.

eq → Pg 26 [Order of B+ Tree] ✗

④ Searching B+ Tree

Time complexity = $\log_p n$

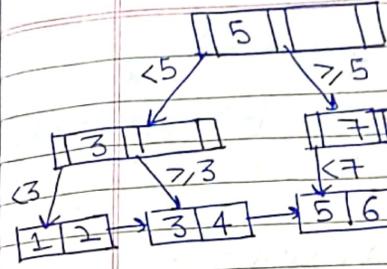
p → much higher than B-Tree ✗

Date _____
Page _____

CLASSEMC
Date _____
Page _____

29

Date _____
Page _____



imp

used for range
Based queries.

→ Search always terminates in leaf in B+ tree not B-Tree. (imp)

eg

→ Pg 28 [Gate - 2015]

[nodes occurred 2.]

* B/B+ Trees are perfectly balanced as distance from root to all leaf nodes are equal. [Gate Q]

eg

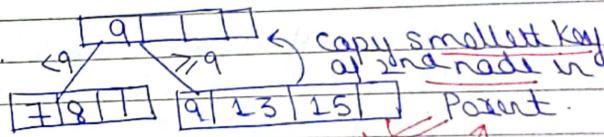
→ Pg 32 [Gate - 2008] [No. of Splits in B-tree insertion] ✗

⑤ B+ Tree Insertion

but always small

Leaf : if order = 5; max Keys = 4.
[7 8 9 13 15]

[7 9 13 15] → insert 8.



Non-leaf

imp

10 20 30 40

Date _____
Page _____ 30

insert 15

20
10 15 30, 40

[10, 15, 20, 30, 40]
[no copy]

eg → Pg 34 [B+ Tree insertion] ✓
Imp. Note: $\lceil \frac{P}{2} \rceil - 1$ in 1 node
& smallest in node is promoted.
Non-Leaf: (7 17 20 25)

$\lceil \frac{P}{2} \rceil - 1$ in 1 node : 17
+ 20 25

* ht of B-tree \geq ht of B+Tree

B+tree grows in breadth
∴ Search time | ∵ only key is present in internal node, ∴ more children.
B+tree: Range based queries easily implemented. ✓

4) B+ Tree Deletion

If any of siblings have keys greater than minimum: Borrow.
If not: merge

In case of underflow
Key $< \lceil \frac{P}{2} \rceil - 1$

Date _____
Page _____ 31

In case of re-arrangements from Sibling:

→ replace the internal node key with the smallest element in right Sub tree.

eg → Pg 36 (Eg 1) [Rearrangements] ✓

eg → Pg 38 [Smallest element in right Subtree is pushed up to satisfy B+tree req] ✓

eg → Pg 39 [Merging] ✓

* PREVIOUS YEARS GATE Q'S

1) SQL

* SQL query can automatically eliminate duplicates.
[RA / TC can].

eg → Pg 131 [Q3] ✓ [R has no dupl. & S is non-empty]

* Select distinct a_1, a_2, \dots, a_n from r_1, r_2, \dots, r_m where P.
 $\sqcup \prod_{a_1, a_2, \dots, a_n} p_i(r_1, r_2, \dots, r_m)$

→ If not a completed SQ, write it & execute the outside P.

eg → Pg 133 [2014] ✓

writing is used while creating Table

eg → Pg 134 [Gate 2014] ✓

eg → Pg 136 [Form examples, see the option & run the SQL query.] ✓

eg → Pg 138 [Gate 2014] (all customers having a good rating) ✓

eg → Pg 139 [Gate - 2016] ✓

* Having clause should be applied only with Group By clause.
All attributes must appear in Select clause (i) in Group by), else put in Aggregate function.

(*) Relational Algebra

→ if $R = m$ tuples, $S = n$ tuples,
max size of join = mn ?
min size of join = 0. ✓

* Sum can't be computed using basic relational Algebra operations.

RA $\xrightarrow{\text{some power}}$ Safe Tuple Calculus

Date
Page

32

CLASSTIME
Date
Page

33

eg → Pg 142 [Gate - 2004] ✓

[Try to solve without eg]

* If common attribute 'a' is not a PK/CK, we get Spurious tuples.
Where tuples in case of loop.

eg → Pg 144 [IT - 2006] ✓

eg → Pg 145 [at least one girl not registered] ✓

eg → Pg 146 [Use eg] ✓

* For a join b/w relation R & relation S, using Nested Loop Method.
These are 3 buffers.

if $\text{size}(R) < \text{size}(S)$,
your disk block access will be:
when: (R) is in the outer loop.

eg → Pg 150 [Gate 2014] [Try to solve by using without eg] ✓

eg → Pg 151 [Gate 2007] ✓
[Supervisor Q].

→ Pg 152 [2008]

→ Pg 153 [Safe & unsafe relational
calculator]. ✓

→ Safe: Tuple from domain.

* Sec ways of not (\neg) .

$\{t \mid \neg(t \in R_i)\} \rightarrow t$ may
belong to many tuples
in Universe : unsafe.

3) TRANSACTION MANAGEMENT

→ Gate 2012 [Pg 154] [Try all
concurrent interleaving of T1 &
T2]. ✓ [just conflicting operations]

→ Pg 155 [Gate 2014]
 $r_1(n), r_2(n), w_1(x), r_3(n), w_2(n)$.
[Do it part]. ✓

→ Pg 157 [Conflict serializable &
recoverable]. ✓

No dirty read: consistent ; &
recoverable Schedule

Here, ta writes data b4 t1 aborts.

→ Pg 158 Gate 2016 ✓

→ Pg 159 [Strict & PLocking] ✓

4) FILE STRUCTURES (Prev)

→ 6.16 (Grp) ↗ imp

→ B+ tree is preferred to BST, since
data transfer from disk is in
Blocks.

→ 6.5 (B-Tree Q) ✓

→ B+ Tree (Insertion & deletion),
Q 6.12, 6.13 → imp Q ✓

→ Q 6.15 [Secondary indexing] ✓

* TEST ANALYSIS

* If asked CK, find closure of each
option and check.

* For dependency preserving, find ant closure,
then project, and see if you can
derive original FDs.

* If not conflict Serializable, Schedule
must be allowed by JPL

→ Q 5 ✓ imp

Vimp

classmate

Date _____

Page 36

* No. of Block transfers in Nested Loop
Join Method =

$$[B_{\text{outer}} + N_{\text{outer}} \times B_{\text{inner}}] \text{ (tuple)} \quad [B_{\text{outer}} \times B_{\text{inner}}] \text{ (block)}$$

* Outer relation must be smaller

* Previous Year Gate Qs

1) SQL

* An SQL query will work well, if there are no indexes on relations.

[SQL doesn't permit attribute names to be repeated in same relation]

eg → Q3 (Pg 131) [r has no duplicates & S is non-empty]
 if duplicates, result ≠ 1.
 (from $r_1, r_2, \dots = r_1 \times r_2 \times r_3$).
 where = Selection (σ)
 Select = projection (Π)

* Duplicates only used while creating Tables.

eg → Pg 134 [Gate 2005]

eg → Pg 136 [Whenever they use raw sets, take distinct values, not duplicates]

classmate

Vimp

Date _____

Page _____

eg → Pg 137 [Gate 2007] ✓
 [correlated SQL, not exists].
 [Take eg to answer this Q].

⇒ Natural Join → atleast one not
 ↗ (atleast one not) = all.

↗ (atleast one customer hasn't given
 good rating) = all customers gave
 good rating.

* [Always take small example].

eg → Pg 139 [Gate 2016] ✗

* Having clause should be applied only
 with Group By clause.
 All attributes must appear in Select
 clause if in group by. Vimp

* Relational algebra has same power
 as SQL relational calculus.

2) Relational Algebra

eg → Pg 142 [Gate 2004] ✗
 Solve first by reading expression.

* If $R(A, B, C, D)$ is decomposed into
 $R_1(A, B, C)$; $R_2(A, D)$;
 then joined to form $S(A, B, C, D)$ [NJ].
 i). $R_1 \cap R_2 = A$ (not C); [Many]
 then, S will contain spurious tuples.
 ∴ $[R \cap S]$ ✗

$\pi_{\text{studId}}(\pi_{\text{courseId}}(R))$ (all female students) \rightarrow
 studId of female all female students
 Date Page
 all the female students paired with all courses.
 Every female student is paired with every course.

eq Pg 145 [course in which only subset of female students have enrolled, not all] ✓

eq Pg 146 [Good Q on NJ] (column matched once).
 R S
 $\begin{array}{ccccc} p & q & R_1 & R_2 & R_3 \\ 1 & 2 & a & b & c \\ 2 & 3 & d & e & f \end{array}$ $\begin{array}{ccccc} p & q & S_1 & S_2 \\ 1 & 2 & a & b \\ 2 & 1 & d & f \end{array}$

$$\therefore \pi_{\text{courseId}}(\pi_{\text{studId}}(R)) \bowtie \pi_{\text{courseId}}(\pi_{\text{studId}}(S))$$

$$\Rightarrow \begin{array}{ccccc} 1 & 1 & & & \vdots \\ 2 & 2 & & & \end{array} \quad \vdots \quad \begin{array}{c} p \\ 1 \\ 2 \end{array} \quad \checkmark$$

eq Pg 151 [Gate 2007] (ump)
 [Correlated Subquery using Tuple Calculus] ✗

a) $\{t | \exists u \in R, t[A] = u[A] \wedge \exists s \in R_a (t[A] = s[A])\}$
 return tuple such that there exists same u in R, such that $t[A] = u[A]$ and there exists no s belonging to R_a , where $t[A] = s[A]$.
 o/p comes from R_1 (SQL). ✗

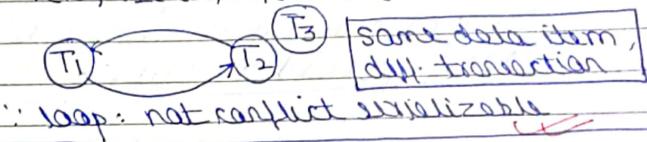
classmate Date Page 39
 $\{t | \neg(t \in R_1)\}$.
 tuple t, such that t does not belong to R_1 . (implies) ✗

3) TRANSACTION MANAGEMENT.

* T_1 T_2 $\xrightarrow{\text{Vimp}}$
 $\text{read}(P)$ $\text{read}(Q)$ $\text{read}(P)$
 $\text{if } P=0, \text{then } Q++$ $\text{if } Q=0, \text{then } P++$
 $\text{write}(Q)$ $\text{write}(P)$

\Rightarrow non-serial interleaving of T_1, T_2 .
 [consider all case of conflicting opn's, and draw Precedence Graphs]

eq Pg 165 ✗
 (A) $r_1(x); r_2(x); w_1(x); r_3(x); w_2(x)$ $\xrightarrow{\text{Schedule}}$



$T_1 \quad T_2 \quad T_3 \quad T_4$
 $r(x)$ $w(x)$ $\xrightarrow{\text{non-serial}}$

$w(x)$ com. \therefore [Nondirty read rule].
 $w(y)$ com. \therefore consistency & strict log.

$w(y)$
 $r(z)$
 com.

$r(x)$
 $w(y)$
 comm.

eg → Pg 159 [Strict 2 Phase Locking]
(unlock exclusive locks on data item only after commit) ✓

* TRANSACTION CONTROL

- If a Schedule is not conflict serializable it doesn't have any blind write. **Vimp**
- It can also be view Serializable.
- If not Serializable, can't achieve 2PL Protocol. ✓

* Database **Vimp**

↓
Pages

↓
Table

↓
Rows

↓
As we move downward,

concurrency increase.

: Day of concurrency is highest, when lock applied on rows. ✓

* A non-serializable Schedule contravenes transaction atomicity. **Nimp**

* In Strict 2PL locking, 2 conditions must be met:

(i) write operation needs exclusive lock on data. **Nimp**

(ii) unlock X-locks only after commit. **A5.9** ✓

* To find if 2 schedules are conflict equivalent to each other, draw a list of conflicting operations and match. [If same, conflict equivalent].
→ Q. 5.10 [Soln in Temp formula list]

eg → Q. 5.14 ✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓

✓